

HomePricePrediction

April 11, 2024

1 Toronto Housing Price Prediction Model

1.1 Introduction

This notebook details our efforts to develop a reliable housing price prediction model for the Toronto area. Given the absence of a robust model for this region, we utilized the Zillow API to gather data for training.

1.2 Data Collection

Training data was acquired using the Zillow API (see `get-data/zillow.py` script). The dataset includes various features of homes recently sold, or available for sale in GTA area.

1.3 Preprocessing

The dataset underwent significant preprocessing: - Removal of columns with excessive missing values or low predictive power. - Encoding of categorical variables such as `homeType` to numerical format suitable for modeling. - Removal of duplicate entries and entries with zero price.

1.4 Model Development

Two models were developed and evaluated: 1. **Linear Regression** 2. **Random Forest**

2 Loading the Dataset

```
[ ]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler
      import numpy as np
      import seaborn as sns
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns
      from scipy import stats
      from scipy.stats import norm, skew

      def load_dataset(filename):
          df = pd.read_csv(filename)
```

```

    return df
df = load_dataset("zillow_homes_sold_toronto.csv")

```

```
[ ]: df['homeStatusForHDP'].value_counts()
```

```
[ ]: homeStatusForHDP
RECENTLY_SOLD    5042
FOR_RENT         98
FOR_SALE         80
Name: count, dtype: int64
```

```
[ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5220 entries, 0 to 5219
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   isShowcaseListing                    5220 non-null   bool
1   longitude                            412 non-null    float64
2   timeOnZillow                        5100 non-null   float64
3   zestimate                           54 non-null     float64
4   daysOnZillow                        5220 non-null   int64
5   zipid                               5220 non-null   int64
6   taxAssessedValue                    54 non-null     float64
7   isUnmappable                        5220 non-null   bool
8   priceForHDP                         5220 non-null   int64
9   dateSold                           5220 non-null   int64
10  state                               5220 non-null   object
11  isFeatured                          5220 non-null   bool
12  isPremierBuilder                    5220 non-null   bool
13  isPreforeclosureAuction             5220 non-null   bool
14  lotAreaValue                        1990 non-null   float64
15  isNonOwnerOccupied                  5220 non-null   bool
16  homeStatus                          5220 non-null   object
17  latitude                            412 non-null    float64
18  lotAreaUnit                         1990 non-null   object
19  bedrooms                           5220 non-null   int64
20  zipcode                             5220 non-null   object
21  homeStatusForHDP                    5220 non-null   object
22  isZillowOwned                       5220 non-null   bool
23  shouldHighlight                     5220 non-null   bool
24  homeType                            5220 non-null   object
25  bathrooms                           5220 non-null   float64
26  rentZestimate                       55 non-null     float64
27  price                               5220 non-null   int64
28  city                                5220 non-null   object
29  streetAddress                       5220 non-null   object

```

```

30 country          5220 non-null object
31 currency         5220 non-null object
32 listing_sub_type 5220 non-null object
33 livingArea       1251 non-null float64
dtypes: bool(8), float64(9), int64(6), object(11)
memory usage: 1.1+ MB

```

```
[ ]: print(df.isnull().sum())
```

```

isShowcaseListing      0
longitude              4808
timeOnZillow           120
zestimate              5166
daysOnZillow          0
zpid                   0
taxAssessedValue       5166
isUnmappable           0
priceForHDP            0
dateSold               0
state                  0
isFeatured             0
isPremierBuilder       0
isPreforeclosureAuction 0
lotAreaValue           3230
isNonOwnerOccupied     0
homeStatus             0
latitude               4808
lotAreaUnit            3230
bedrooms               0
zipcode                0
homeStatusForHDP       0
isZillowOwned          0
shouldHighlight        0
homeType               0
bathrooms              0
rentZestimate          5165
price                  0
city                   0
streetAddress          0
country                0
currency               0
listing_sub_type       0
livingArea             3969
dtype: int64

```

3 Data Preprocessing and Clean up

3.1 Removing columns

Columns like state, homeStatus, lotAreaUnit, homeType, city, streetAddress, country, and currency are categorical and should be encoded to numerical values if you plan to use them in your model.

```
[ ]: def drop_columns(df):
    columns_with_many_missing = df.columns[df.isnull().mean() > .99]
    columns_with_low_predictive_power = [
        'zpid', 'streetAddress', 'listing_sub_type',
        'isUnmappable', 'isShowcaseListing', 'isFeatured',
        'isPremierBuilder', 'isPreforeclosureAuction',
        'isNonOwnerOccupied', 'isZillowOwned', 'shouldHighlight',
        'city', 'streetAddress', 'country', 'currency', 'zpid', 'lotAreaUnit',
        'homeStatusForHDP', 'homeStatus', 'state',
        'zestimate', 'taxAssessedValue', 'rentZestimate',
        'zipcode',
        'dateSold',
        'priceForHDP',
        # 'daysOnZillow'
    ]

    columns_to_drop = list(set(columns_with_many_missing.to_list() +
    ↪columns_with_low_predictive_power))

    df_dropped = df.drop(columns=columns_to_drop)

    print(f"Number of rows: {df.shape[0]}")
    print("columns_with_many_missing:", columns_with_many_missing)
    print("Columns dropped:", columns_to_drop)
    print("\nRemaining columns:", df_dropped.columns.tolist())

    df_dropped.head()
    return df_dropped

df = drop_columns(df)
```

Number of rows: 5220

columns_with_many_missing: Index([], dtype='object')

Columns dropped: ['isUnmappable', 'country', 'dateSold', 'rentZestimate',
'isNonOwnerOccupied', 'priceForHDP', 'homeStatus', 'isPreforeclosureAuction',
'zpid', 'zipcode', 'currency', 'city', 'listing_sub_type', 'isFeatured',
'state', 'shouldHighlight', 'isShowcaseListing', 'streetAddress', 'lotAreaUnit',
'taxAssessedValue', 'isPremierBuilder', 'isZillowOwned', 'zestimate',
'homeStatusForHDP']

Remaining columns: ['longitude', 'timeOnZillow', 'daysOnZillow', 'lotAreaValue',
'latitude', 'bedrooms', 'homeType', 'bathrooms', 'price', 'livingArea']

3.2 Categorical Data Encoding

Since `homeType` is a categorical variable, we'll need to encode it to use it in most machine learning models, which require numerical input.

One-Hot Encoding: This method converts the categorical variable into a series of binary columns, each representing a unique category. One-hot encoding is suitable when there is no ordinal relationship between the categories. Since `homeType` likely does not have a natural order, this method is appropriate.

```
[ ]: df['homeType'].value_counts()
```

```
[ ]: homeType
SINGLE_FAMILY    3577
CONDO           905
APARTMENT       461
TOWNHOUSE       222
MULTI_FAMILY     55
Name: count, dtype: int64
```

3.3 Handling Sparse Categories

Given the distribution of your `homeType` categories, it appears that `MANUFACTURED`, `MULTI_FAMILY`, and `TOWNHOUSE` categories are relatively sparse compared to the others. To clean up the `homeType` column and ensure your model remains robust, you could group these less frequent categories into an “Other” category. This approach simplifies your model while retaining the information these categories provide.

```
[ ]: print(df.columns)
# print(type(df))
# df['homeType'] = df['homeType'].replace(['MANUFACTURED', 'MULTI_FAMILY',
# ↪ 'TOWNHOUSE'], 'Other')

# print("\nNew distribution of 'homeType':")
# print(df['homeType'].value_counts())
# print(type(df))
```

```
Index(['longitude', 'timeOnZillow', 'daysOnZillow', 'lotAreaValue', 'latitude',
      'bedrooms', 'homeType', 'bathrooms', 'price', 'livingArea'],
      dtype='object')
```

```
[ ]: def encode_homeType(df):
    from sklearn.preprocessing import OneHotEncoder

    # Initializing the OneHotEncoder without specifying 'sparse'
    encoder = OneHotEncoder(drop='first') # Assuming 'drop' argument is ↪
    ↪ supported in your version

    # Fit and transform 'homeType' column
```

```

homeType_encoded_sparse = encoder.fit_transform(df[['homeType']])

# Convert the sparse matrix to a dense matrix and then to a DataFrame
homeType_encoded_df = pd.DataFrame(homeType_encoded_sparse.toarray(),
↳columns=encoder.get_feature_names_out(['homeType']))

# Concatenate the encoded DataFrame back with the original DataFrame
↳(excluding the original 'homeType' column)
df_final = pd.concat([df.drop('homeType', axis=1), homeType_encoded_df],
↳axis=1)

print(f"Number of rows: {df_final.shape[0]}")
df_final.head()
return df_final

df = encode_homeType(df)

```

Number of rows: 5220

3.4 Removing duplicates

```

[ ]: def remove_duplicates(df):
    # Check for duplicate rows
    print("Duplicates before:", df.duplicated().sum())

    # Remove duplicate rows
    df_removed = df.drop_duplicates()

    # Check again to confirm removal
    print("Duplicates after:", df_removed.duplicated().sum())

    print("Number of rows:", len(df_removed))
    df_removed.head()
    return df_removed

df = remove_duplicates(df)

```

Duplicates before: 857

Duplicates after: 0

Number of rows: 4363

3.5 Remove where price is 0

```

[ ]: # print((df['price'] == 0).sum())
    # print((df['priceForHDP'] == 0).sum())
    # print(((df['price'] == 0) & (df['priceForHDP'] != 0)).sum())
    # df = df[df['price'] != 0]
    # print('-----')

```

```
# print((df['price'] == 0).sum())
# print((df['priceForHDP'] == 0).sum())
# print(((df['price'] == 0) & (df['priceForHDP'] != 0)).sum())

# num_rows = df.shape[0]
# print(f"Number of rows: {num_rows}")
```

4 Exploratory Data Analysis (EDA)

4.1 Describe

```
[ ]: # display basic descriptive statistics for numerical columns
print(df.describe())
print(f"Number of rows: {df.shape[0]}")

df.head()
```

| | longitude | timeOnZillow | daysOnZillow | lotAreaValue | latitude \ |
|-------|------------|--------------|--------------|--------------|------------|
| count | 355.000000 | 4.357000e+03 | 4363.000000 | 1496.000000 | 355.000000 |
| mean | -79.400171 | 6.187088e+10 | 714.909008 | 3913.129733 | 43.919090 |
| std | 0.714517 | 1.791425e+10 | 208.894836 | 2488.899417 | 0.437308 |
| min | -81.625465 | 2.004717e+09 | -1.000000 | 0.000000 | 42.968273 |
| 25% | -79.784454 | 5.298072e+10 | 609.000000 | 2330.000000 | 43.667984 |
| 50% | -79.419914 | 6.222551e+10 | 720.000000 | 3234.000000 | 43.777637 |
| 75% | -79.351850 | 7.458072e+10 | 859.500000 | 5445.000000 | 44.275513 |
| max | -77.059010 | 9.462557e+10 | 1095.000000 | 10419.000000 | 44.741947 |

| | bedrooms | bathrooms | price | livingArea | homeType_CONDO \ |
|-------|-------------|-------------|--------------|-------------|------------------|
| count | 4363.000000 | 4363.000000 | 4.363000e+03 | 908.000000 | 4363.000000 |
| mean | 3.316755 | 2.675338 | 1.402629e+06 | 2190.943833 | 0.203071 |
| std | 1.664849 | 1.339355 | 8.614329e+05 | 1090.422818 | 0.402331 |
| min | 0.000000 | 0.000000 | 1.600000e+03 | 368.000000 | 0.000000 |
| 25% | 2.000000 | 2.000000 | 8.311000e+05 | 1471.000000 | 0.000000 |
| 50% | 3.000000 | 2.000000 | 1.250000e+06 | 1980.000000 | 0.000000 |
| 75% | 5.000000 | 4.000000 | 1.650000e+06 | 2825.000000 | 0.000000 |
| max | 8.000000 | 8.000000 | 5.200000e+06 | 6534.000000 | 1.000000 |

| | homeType_MULTI_FAMILY | homeType_SINGLE_FAMILY | homeType_TOWNHOUSE |
|-------|-----------------------|------------------------|--------------------|
| count | 4363.000000 | 4363.000000 | 4363.000000 |
| mean | 0.012606 | 0.627779 | 0.050882 |
| std | 0.111579 | 0.483452 | 0.219783 |
| min | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.000000 | 0.000000 | 0.000000 |
| 50% | 0.000000 | 1.000000 | 0.000000 |
| 75% | 0.000000 | 1.000000 | 0.000000 |
| max | 1.000000 | 1.000000 | 1.000000 |

Number of rows: 4363

```
[ ]: longitude  timeOnZillow  daysOnZillow  lotAreaValue  latitude  bedrooms  \
0          NaN  6.205267e+10             718          NaN          NaN          0
1 -79.419914  8.036947e+10             930          NaN  43.661144          0
2          NaN  8.131987e+10             941          NaN          NaN          0
3          NaN  8.563987e+10             991          NaN          NaN          0
4          NaN  6.205267e+10             718          NaN          NaN          0

    bathrooms    price  livingArea  homeType_CONDO  homeType_MULTI_FAMILY  \
0          1.0   430000          NaN             1.0                     0.0
1          1.0  1452000        750.0             0.0                     0.0
2          1.0   430000          NaN             1.0                     0.0
3          4.0  1300000          NaN             0.0                     0.0
4          1.0   430000          NaN             1.0                     0.0


    homeType_SINGLE_FAMILY  homeType_TOWNHOUSE
0                      0.0                   0.0
1                      1.0                   0.0
2                      0.0                   0.0
3                      1.0                   0.0
4                      0.0                   0.0
```

4.2 Coorelation Matrix

```
[ ]: # this function is adopted from https://www.kaggle.com/code/shtrausslearning/
    ↪ bayesian-regression-house-price-prediction/notebook

def corrMat(df,id=False):
    corr_mat = df.corr().round(2)
    f, ax = plt.subplots(figsize=(8,8))
    # mask = np.zeros_like(corr_mat,dtype=np.bool)
    mask = np.zeros_like(corr_mat, dtype=bool)
    mask[np.triu_indices_from(mask)] = True
    sns.heatmap(corr_mat,mask=mask,vmin=-1,vmax=1,center=0,
                cmap='plasma',square=False,lw=2,annot=True,cbar=False);plt.
    ↪ show()

# X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)
corrMat(df)
```

bathrooms have a strong positive correlation with bedrooms (0.7), which is expected since more bedrooms often mean more bathrooms. based on this we create a ‘Total Rooms’ Feature: Since bathrooms and bedrooms are highly correlated, we create a new feature that combines them, such as $\text{total_rooms} = \text{bathrooms} + \text{bedrooms}$. This could help the model by consolidating two highly correlated features into one.

4.3 Add bedrooms+bathrooms

```
[ ]: def combine_and_drop(df):
    df['total_rooms'] = df['bedrooms'] + df['bathrooms']
    df = df.drop(['bedrooms', 'bathrooms'], axis=1)
    df.head()
    return df
```

```
df = combine_and_drop(df)
```

4.4 remove nan

```
[ ]: def print_categorical_value_counts(df):
    categorical_columns = df.select_dtypes(include=['object', 'bool']).columns

    print("Categorical columns:" + str(categorical_columns))

    for column in categorical_columns:
        print(f"Value counts for {column}:")
        print(df[column].value_counts())
        print("\n")

    print(f"Number of rows: {df.shape[0]}")
    df.head()
    return df

df = print_categorical_value_counts(df)
```

Categorical columns: Index([], dtype='object')

Number of rows: 4363

```
[ ]: print(df.columns[df.isna().any()].tolist())

nan_cols = df.columns[df.isna().any()].tolist()

# df = df[df['price'].notna()]
# df = df[df['bedrooms'].notna() & df['bathrooms'].notna()]
# df = df[df['livingArea'].notna()]

for col in nan_cols:
    print(f"{col}: {df[col].isna().sum()}")

# Count columns with NaN values
nan_cols_count = df.columns[df.isna().any()].shape[0]
print(f"Number of columns with NaN values: {nan_cols_count}")

# Remove columns with NaN values
# df = df.dropna(axis=1)
print(f"Number of columns with NaN values: {nan_cols_count}")

print(f"Number of rows: {df.shape[0]}")
# print((df['bedrooms'] == 0).sum())
# df = df[df['bedrooms'] != 0]
```

```
print(f"Number of rows: {df.shape[0]}")
df.head()
```

```
['longitude', 'timeOnZillow', 'lotAreaValue', 'latitude', 'livingArea']
longitude: 4008
timeOnZillow: 6
lotAreaValue: 2867
latitude: 4008
livingArea: 3455
Number of columns with NaN values: 5
Number of columns with NaN values: 5
Number of rows: 4363
Number of rows: 4363
```

```
[ ]:  longitude  timeOnZillow  daysOnZillow  lotAreaValue  latitude  price  \
0         NaN  6.205267e+10          718          NaN         NaN  430000
1 -79.419914  8.036947e+10          930          NaN  43.661144  1452000
2         NaN  8.131987e+10          941          NaN         NaN  430000
3         NaN  8.563987e+10          991          NaN         NaN  1300000
4         NaN  6.205267e+10          718          NaN         NaN  430000

    livingArea  homeType_CONDO  homeType_MULTI_FAMILY  homeType_SINGLE_FAMILY  \
0         NaN              1.0              0.0              0.0
1       750.0              0.0              0.0              1.0
2         NaN              1.0              0.0              0.0
3         NaN              0.0              0.0              1.0
4         NaN              1.0              0.0              0.0

    homeType_TOWNHOUSE  total_rooms
0              0.0              1.0
1              0.0              1.0
2              0.0              1.0
3              0.0              4.0
4              0.0              1.0
```

4.5 Handle missing values

```
[ ]: def handle_missing_values(df):
      df = df[df['price'].notna()]
      df = df.dropna(axis=1)
      return df

df = handle_missing_values(df)
```

5 Training the Model

Implementing a Linear Regression model and split the data into training and testing sets

```
[ ]: # Final columns before training
print(f"Number of rows: {df.shape[0]}")

df.head()
# df.to_csv('training_dataset_saved.csv', index=False)
```

Number of rows: 4363

```
[ ]:  daysOnZillow    price  homeType_CONDO  homeType_MULTI_FAMILY \
0         718    430000             1.0             0.0
1         930   1452000             0.0             0.0
2         941    430000             1.0             0.0
3         991   1300000             0.0             0.0
4         718    430000             1.0             0.0

      homeType_SINGLE_FAMILY  homeType_TOWNHOUSE  total_rooms
0                      0.0             0.0             1.0
1                      1.0             0.0             1.0
2                      0.0             0.0             1.0
3                      1.0             0.0             4.0
4                      0.0             0.0             1.0
```

5.1 Linear Regression

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler
import numpy as np

def train_linear_regression(df, target):
    X = df.drop(target, axis=1)
    y = df[target]

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
    scaler = StandardScaler()

    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    linear_model = LinearRegression()
    linear_model.fit(X_train_scaled, y_train)

    y_pred = linear_model.predict(X_test_scaled)
    return y_test, y_pred

y_test_lin, y_pred_lin = train_linear_regression(df, 'price')
```

5.1.1 Evaluation

```
[ ]: from sklearn.metrics import r2_score

def calculate_metrics(y_test, y_pred, model_name):
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    r_squared = r2_score(y_test, y_pred)
    print(f"Root Mean Squared Error (RMSE) on Test Set for {model_name}:↳
↳{rmse}")
    print(f"R-squared (R2) on Test Set for {model_name}: {r_squared}")
    ↳
    ↳print("-----\n")

calculate_metrics(y_test_lin, y_pred_lin, 'Linear Regression')
```

Root Mean Squared Error (RMSE) on Test Set for Linear Regression:

704751.6957342114

R-squared (R2) on Test Set for Linear Regression: 0.2774396916125408

5.2 Random Forest

```
[ ]: from sklearn.ensemble import RandomForestRegressor

def train_random_forest(df, target):
    X = df.drop(target, axis=1)
    y = df[target]

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
    scaler = StandardScaler()

    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    random_forest_model = RandomForestRegressor(n_estimators=100,
↳random_state=42)
    random_forest_model.fit(X_train_scaled, y_train)

    y_pred = random_forest_model.predict(X_test_scaled)
    return y_test, y_pred
y_test_rf, y_pred_rf = train_random_forest(df, 'price')
```

5.2.1 Evaluation

```
[ ]: calculate_metrics(y_test_rf, y_pred_rf, "Random Forest")
```

Root Mean Squared Error (RMSE) on Test Set for Random Forest: 155818.39580078336
R-squared (R2) on Test Set for Random Forest: 0.9646784868734193

5.3 Update coordinates and rerun

We cannot use addresses directly in the model, and the original dataset had only a few records with both latitude and longitude set. However, addresses are crucial for home price prediction. To incorporate addresses into the model, we used a Python script (get-data/geoapify.py) to call the Geoapify API and fetch coordinates (latitude and longitude) for the addresses. These coordinates were then saved in a CSV file (zillow_homes_sold_toronto_coordinates_added.csv), and the dataset was updated with the retrieved coordinates. The models will be rerun using this new dataset.

```
[ ]: df = load_dataset("zillow_homes_sold_toronto_cordinations_added.csv")
df = drop_columns(df)
df = encode_homeType(df)
df = remove_duplicates(df)
df = combine_and_drop(df)
df = handle_missing_values(df)
```

```
df.head()
```

Number of rows: 5220

columns_with_many_missing: Index([], dtype='object')

Columns dropped: ['isUnmappable', 'country', 'dateSold', 'rentZestimate', 'isNonOwnerOccupied', 'priceForHDP', 'homeStatus', 'isPreforeclosureAuction', 'zpid', 'zipcode', 'currency', 'city', 'listing_sub_type', 'isFeatured', 'state', 'shouldHighlight', 'isShowcaseListing', 'streetAddress', 'lotAreaUnit', 'taxAssessedValue', 'isPremierBuilder', 'isZillowOwned', 'zestimate', 'homeStatusForHDP']

Remaining columns: ['longitude', 'timeOnZillow', 'daysOnZillow', 'lotAreaValue', 'latitude', 'bedrooms', 'homeType', 'bathrooms', 'price', 'livingArea']

Number of rows: 5220

Duplicates before: 857

Duplicates after: 0

Number of rows: 4363

```
[ ]:  longitude  daysOnZillow  latitude  price  homeType_CONDO  \
0 -79.400083         718  43.636951   430000           1.0
1 -79.419914         930  43.661144  1452000           0.0
2 -79.400083         941  43.636951   430000           1.0
3 -80.344809         991  43.355007  1300000           0.0
4 -79.400083         718  43.636951   430000           1.0

      homeType_MULTI_FAMILY  homeType_SINGLE_FAMILY  homeType_TOWNHOUSE  \
0                0.0                0.0                0.0
1                0.0                1.0                0.0
2                0.0                0.0                0.0
3                0.0                1.0                0.0
4                0.0                0.0                0.0

      total_rooms
0                1.0
1                1.0
2                1.0
3                4.0
4                1.0
```

```
[ ]: y_test_lin, y_pred_lin = train_linear_regression(df, 'price')
      calculate_metrics(y_test_lin, y_pred_lin, "Linear Regression")

      y_test_rf, y_pred_rf = train_random_forest(df, 'price')
      calculate_metrics(y_test_rf, y_pred_rf, "Random Forest")
```

Root Mean Squared Error (RMSE) on Test Set for Linear Regression:

666110.8385600615

R-squared (R2) on Test Set for Linear Regression: 0.35450208646839176

Root Mean Squared Error (RMSE) on Test Set for Random Forest: 74350.92356764115
R-squared (R2) on Test Set for Random Forest: 0.991957805405347

6 RMSE Analysis

The RMSE of 74,350.92 is considered in the context of the dataset's statistics:

Average Price Context: The average price in the dataset is approximately \$1,402,629. An RMSE of 74,350.92 represents about 5.3% of the average price, indicating a relatively small error relative to the average home value.

Compared to Standard Deviation: The standard deviation of the prices is \$861,432.95, suggesting significant variability in house prices. An RMSE much smaller than the standard deviation indicates that the model predictions are reasonably accurate within the existing price variability.

Relative Error Scale: Considering the price range in the dataset, from a minimum of \$1,600 to a maximum of \$5,200,000, an RMSE of 74,350.92 demonstrates a moderate level of precision, especially effective for middle and high-end properties.

7 Conclusion

The Random Forest model significantly outperformed the Linear Regression model, indicating its better suitability for this application. The results validate the effectiveness of the Random Forest model in predicting housing prices with high accuracy. Future efforts might focus on integrating more dynamic features and improving data collection to refine predictions further. This is likely due to:

Non-linearity: Random Forest can capture complex, non-linear relationships between features that affect housing prices

Robust to Outliers: This model is less sensitive to extreme values, which are common in real estate markets, like the occasional luxury home sales that do not fit typical pricing patterns.

Feature Interactions: Random Forest automatically considers how combinations of features impact house prices. For example, the model can intuitively assess how the interaction between property size and its location (represented by latitude and longitude) affects its value, unlike Linear Regression, which requires explicitly crafted interaction terms to capture such complex relationships.

Prevention of Overfitting: Random Forest uses techniques like bootstrapping and feature randomness to build diverse trees that generalize better to new data, reducing the risk of overfitting seen in complex Linear Regression models.

Handles Various Data Types: It effortlessly processes different data types and does not require extensive preprocessing for categorical variables, streamlining the modeling process.

8 Model improvements

Following potentially can be done to improve model's performance:

Feature Engineering: Develop new features that capture more nuances of the real estate market, such as proximity to amenities or property condition or walking score.

Data Enrichment: Integrate additional data sources, such as economic indicators or demographic statistics, to provide more context for predictions.

Advanced Modeling Techniques: Experiment with more complex models like Gradient Boosting Machines (GBM) or neural networks that might capture patterns better.

Parameter Tuning: Optimize hyperparameters of the Random Forest model through grid search or random search to enhance model accuracy.

Handling Outliers: Identify and manage outliers in the dataset that may skew the model predictions.